

Cost-Sensitive Label Embedding for Multi-Label Classification

Kuan-Hao Huang

Advisor: Hsuan-Tien Lin

Department of Computer Science & Information Engineering
National Taiwan University



TAAI, November 27, 2016

Multi-Label Classification (MLC)

Multi-label classification





- ▶ an **extension** of the **multiclass classification**
- ▶ allow instance with **multiple** associated classes

Multi-Label Classification (MLC)

Multi-label classification

- ▶ an **extension** of the **multiclass classification**
- ▶ allow instance with **multiple** associated classes

Example: image tag with (dog, cat, rabbit, shark)

image				
tag	{ dog, cat }	{ dog }	{ dog, cat, rabbit }	{ shark }
label	(1, 1, 0, 0)	(1, 0, 0, 0)	(1, 1, 1, 0)	(0, 0, 0, 1)

Multi-Label Classification (MLC)

Notation

- ▶ feature vector (image): $\mathbf{x} \in \mathcal{X} \subseteq \mathbb{R}^d$
- ▶ label vector (tag): $\mathbf{y} \in \mathcal{Y} \subseteq \{0, 1\}^K$

Multi-Label Classification (MLC)

Notation

- ▶ feature vector (image): $\mathbf{x} \in \mathcal{X} \subseteq \mathbb{R}^d$
- ▶ label vector (tag): $\mathbf{y} \in \mathcal{Y} \subseteq \{0, 1\}^K$

Multi-label classification (MLC)

- ▶ given training instances $\mathcal{D} = \{(\mathbf{x}^{(n)}, \mathbf{y}^{(n)})\}_{n=1}^N$
- ▶ learn a **predictor** h from \mathcal{D}
- ▶ for testing instance (\mathbf{x}, \mathbf{y}) , prediction $\tilde{\mathbf{y}} = h(\mathbf{x})$
- ▶ let the prediction $\tilde{\mathbf{y}}$ is close to ground truth \mathbf{y}

Multi-Label Classification (MLC)

Notation

- ▶ feature vector (image): $\mathbf{x} \in \mathcal{X} \subseteq \mathbb{R}^d$
- ▶ label vector (tag): $\mathbf{y} \in \mathcal{Y} \subseteq \{0, 1\}^K$

Multi-label classification (MLC)

- ▶ given training instances $\mathcal{D} = \{(\mathbf{x}^{(n)}, \mathbf{y}^{(n)})\}_{n=1}^N$
- ▶ learn a **predictor** h from \mathcal{D}
- ▶ for testing instance (\mathbf{x}, \mathbf{y}) , prediction $\tilde{\mathbf{y}} = h(\mathbf{x})$
- ▶ let the prediction $\tilde{\mathbf{y}}$ is close to ground truth \mathbf{y}

Evaluation of closeness

- ▶ **cost function** $c(\mathbf{y}, \tilde{\mathbf{y}})$: the penalty of predicting \mathbf{y} as $\tilde{\mathbf{y}}$
- ▶ Hamming loss, 0/1 loss, Rank loss, F1 score(loss), Accuracy score(loss)

Cost-Sensitive Multi-Label Classification (CSMLC)

Notation

- ▶ feature vector (image): $\mathbf{x} \in \mathcal{X} \subseteq \mathbb{R}^d$
- ▶ label vector (tag): $\mathbf{y} \in \mathcal{Y} \subseteq \{0, 1\}^K$


Cost-sensitive multi-label classification (CSMLC)

- ▶ given training instances $\mathcal{D} = \{(\mathbf{x}^{(n)}, \mathbf{y}^{(n)})\}_{n=1}^N$ and **cost function** c
- ▶ learn a **predictor** h from **both** \mathcal{D} **and** c
- ▶ for testing instance (\mathbf{x}, \mathbf{y}) , prediction $\tilde{\mathbf{y}} = h(\mathbf{x})$
- ▶ let the prediction $\tilde{\mathbf{y}}$ is close to ground truth \mathbf{y}

Evaluation of closeness

- ▶ **cost function** $c(\mathbf{y}, \tilde{\mathbf{y}})$: the penalty of predicting \mathbf{y} as $\tilde{\mathbf{y}}$
- ▶ Hamming loss, 0/1 loss, Rank loss, F1 score(loss), Accuracy score(loss)


Label Embedding



Training stage

- ▶ **embedding function Φ** : label vector $\mathbf{y} \rightarrow$ embedded vector \mathbf{z}
- ▶ train a regressor g from $\{(\mathbf{x}^{(n)}, \mathbf{z}^{(n)})\}_{n=1}^N$

Label Embedding




Training stage

- ▶ **embedding function Φ** : label vector $\mathbf{y} \rightarrow$ embedded vector \mathbf{z}
- ▶ train a regressor g from $\{(\mathbf{x}^{(n)}, \mathbf{z}^{(n)})\}_{n=1}^N$

Predicting stage

- ▶ for testing instance \mathbf{x} , predicted embedded vector $\tilde{\mathbf{z}} = g(\mathbf{x})$
- ▶ **decoding function Ψ** : predicted embedded vector $\tilde{\mathbf{z}} \rightarrow$ predicted label vector $\tilde{\mathbf{y}}$


Cost-Sensitive Label Embedding



Existing works

- **label embedding**: PLST, FaIE, RA_kEL, ECC-based [Tai et al., 2012; Lin et al., 2014; Tsoumakas et al., 2011; Ferng et al., 2013]
- **cost-sensitivity**: CFT, PCC [Li et al., 2014; Dembczynski et al., 2010]
- **cost-sensitivity + label embedding**: no existing works

Cost-Sensitive Label Embedding




Existing works

- ▶ **label embedding**: PLST, FaIE, RA_kEL, ECC-based [Tai et al., 2012; Lin et al., 2014; Tsoumakas et al., 2011; Ferng et al., 2013]
- ▶ **cost-sensitivity**: CFT, PCC [Li et al., 2014; Dembczynski et al., 2010]
- ▶ **cost-sensitivity + label embedding**: no existing works

Cost-sensitive label embedding

- ▶ consider **cost function** c when designing **embedding function** Φ and **decoding function** Ψ (cost-sensitive embedded vectors z)

Cost-Sensitive Label Embedding (Training)




embedded space \mathcal{Z}

Training stage

- ▶ distances between embedded vectors \Leftrightarrow cost information
- ▶ larger (smaller) distance $d(\mathbf{z}_i, \mathbf{z}_j) \Leftrightarrow$ higher (lower) cost $c(\mathbf{y}_i, \mathbf{y}_j)$

Cost-Sensitive Label Embedding (Training)




embedded space \mathcal{Z}

Training stage

- ▶ distances between embedded vectors \Leftrightarrow cost information
- ▶ larger (smaller) distance $d(\mathbf{z}_i, \mathbf{z}_j) \Leftrightarrow$ higher (lower) cost $c(\mathbf{y}_i, \mathbf{y}_j)$
- ▶ $d(\mathbf{z}_i, \mathbf{z}_j) \approx \sqrt{c(\mathbf{y}_i, \mathbf{y}_j)}$


Cost-Sensitive Label Embedding (Training)



Training stage

- ▶ distances between embedded vectors \Leftrightarrow cost information
- ▶ larger (smaller) distance $d(\mathbf{z}_i, \mathbf{z}_j) \Leftrightarrow$ higher (lower) cost $c(\mathbf{y}_i, \mathbf{y}_j)$
- ▶ $d(\mathbf{z}_i, \mathbf{z}_j) \approx \sqrt{c(\mathbf{y}_i, \mathbf{y}_j)}$


Cost-Sensitive Label Embedding (Predicting)



Predicting stage

- ▶ for testing instance \mathbf{x} , predicted embedded vector $\tilde{\mathbf{z}} = g(\mathbf{x})$

Cost-Sensitive Label Embedding (Predicting)



Predicting stage

- ▶ for testing instance \mathbf{x} , predicted embedded vector $\tilde{\mathbf{z}} = g(\mathbf{x})$
- ▶ find **nearest embedded vector** \mathbf{z}_q of $\tilde{\mathbf{z}}$
- ▶ cost-sensitive prediction $\tilde{\mathbf{y}} = \mathbf{y}_q$

Cost-Sensitive Label Embedding (Theorem)

Theorem

$$c(\mathbf{y}, \tilde{\mathbf{y}}) \leq 5 \left(\underbrace{\left(d(\mathbf{z}, \mathbf{z}_q) - \sqrt{c(\mathbf{y}, \tilde{\mathbf{y}})} \right)^2}_{\text{embedding error}} + \underbrace{d(\mathbf{z}, \tilde{\mathbf{z}})^2}_{\text{regression error}} \right)$$

Cost-Sensitive Label Embedding (Theorem)

Theorem

$$c(\mathbf{y}, \tilde{\mathbf{y}}) \leq 5 \left(\underbrace{(d(\mathbf{z}, \mathbf{z}_q) - \sqrt{c(\mathbf{y}, \tilde{\mathbf{y}})})^2}_{\text{embedding error}} + \underbrace{d(\mathbf{z}, \tilde{\mathbf{z}})^2}_{\text{regression error}} \right)$$

Optimization

- ▶ **embedding error** \rightarrow multidimensional scaling (manifold learning)
- ▶ **regression error** \rightarrow regressor g

Cost-Sensitive Label Embedding (Theorem)

Theorem

$$c(\mathbf{y}, \tilde{\mathbf{y}}) \leq 5 \left(\underbrace{\left(d(\mathbf{z}, \mathbf{z}_q) - \sqrt{c(\mathbf{y}, \tilde{\mathbf{y}})} \right)^2}_{\text{embedding error}} + \underbrace{d(\mathbf{z}, \tilde{\mathbf{z}})^2}_{\text{regression error}} \right)$$


Optimization

- ▶ **embedding error** \rightarrow multidimensional scaling (manifold learning)
- ▶ **regression error** \rightarrow regressor g

Challenge

- ▶ **asymmetric cost function vs. symmetric distance?**
- ▶ $c(\mathbf{y}_i, \mathbf{y}_j) \neq c(\mathbf{y}_j, \mathbf{y}_i)$ vs. $d(\mathbf{z}_i, \mathbf{z}_j)$


Mirroring Trick



embedded space \mathcal{Z}

- ▶ two roles of y : **ground truth role** $z_i^{(t)}$ and **prediction role** $z_i^{(p)}$


Mirroring Trick



embedded space \mathcal{Z}

- ▶ two roles of y : **ground truth role** $z_i^{(t)}$ and **prediction role** $z_i^{(p)}$
- ▶ predict y_i as $y_j \Rightarrow \sqrt{c(y_i, y_j)}$ for $z_i^{(t)}$ and $z_j^{(p)}$


Mirroring Trick



embedded space \mathcal{Z}


- ▶ two roles of y : **ground truth role** $z_i^{(t)}$ and **prediction role** $z_i^{(p)}$
- ▶ predict y_i as $y_j \Rightarrow \sqrt{c(y_i, y_j)}$ for $z_i^{(t)}$ and $z_j^{(p)}$
- ▶ predict y_j as $y_i \Rightarrow \sqrt{c(y_j, y_i)}$ for $z_i^{(p)}$ and $z_j^{(t)}$

Mirroring Trick



- ▶ two roles of \mathbf{y} : **ground truth role** $\mathbf{z}_i^{(t)}$ and **prediction role** $\mathbf{z}_i^{(p)}$
- ▶ predict \mathbf{y}_i as $\mathbf{y}_j \Rightarrow \sqrt{c(\mathbf{y}_i, \mathbf{y}_j)}$ for $\mathbf{z}_i^{(t)}$ and $\mathbf{z}_j^{(p)}$
- ▶ predict \mathbf{y}_j as $\mathbf{y}_i \Rightarrow \sqrt{c(\mathbf{y}_j, \mathbf{y}_i)}$ for $\mathbf{z}_i^{(p)}$ and $\mathbf{z}_j^{(t)}$
- ▶ learn **regressor** g from $\mathbf{z}_i^{(p)}, \mathbf{z}_2^{(p)}, \dots, \mathbf{z}_L^{(p)}$

Mirroring Trick



- ▶ two roles of \mathbf{y} : **ground truth role** $\mathbf{z}_i^{(t)}$ and **prediction role** $\mathbf{z}_i^{(p)}$
- ▶ predict \mathbf{y}_i as $\mathbf{y}_j \Rightarrow \sqrt{c(\mathbf{y}_i, \mathbf{y}_j)}$ for $\mathbf{z}_i^{(t)}$ and $\mathbf{z}_j^{(p)}$
- ▶ predict \mathbf{y}_j as $\mathbf{y}_i \Rightarrow \sqrt{c(\mathbf{y}_j, \mathbf{y}_i)}$ for $\mathbf{z}_i^{(p)}$ and $\mathbf{z}_j^{(t)}$
- ▶ learn **regressor** g from $\mathbf{z}_i^{(p)}, \mathbf{z}_2^{(p)}, \dots, \mathbf{z}_L^{(p)}$
- ▶ find **nearest embedded vector** of $\tilde{\mathbf{z}}$ from $\mathbf{z}_1^{(t)}, \mathbf{z}_2^{(t)}, \dots, \mathbf{z}_L^{(t)}$

Candidate Set

Challenge

- ▶ label vector $\mathbf{y} \in \mathcal{Y} \subseteq \{0, 1\}^K$
- ▶ 2^K possible label vectors (too many)
- ▶ what is the important(useful) label vectors?

Candidate Set

Challenge

- ▶ label vector $\mathbf{y} \in \mathcal{Y} \subseteq \{0, 1\}^K$
- ▶ 2^K possible label vectors (**too many**)
- ▶ what is the important(useful) label vectors?

Candidate Set

- ▶ consider a **candidate set** \mathcal{S} instead of \mathcal{Y}
- ▶ only label vectors in \mathcal{S} are embedded
- ▶ \mathcal{S}_{train} (all the label vectors in training set) is a reasonable choice

Cost-Sensitive Label Embedding with Multidimensional Scaling

Training stage of **CLEMS**

- ▶ given training instances $\mathcal{D} = \{(\mathbf{x}^{(n)}, \mathbf{y}^{(n)})\}_{n=1}^N$ and cost function c
- ▶ determine the candidate set \mathcal{S}
- ▶ find $\mathbf{z}_i^{(t)}$ and $\mathbf{z}_i^{(p)}$ for all $\mathbf{y}_i \in \mathcal{S}$ by **multidimensional scaling**
- ▶ $\Phi: \mathbf{y}_i \rightarrow \mathbf{z}_i^{(p)}$
- ▶ train a multi-target regressor g from $\{(\mathbf{x}^{(n)}, \Phi(\mathbf{y}^{(n)}))\}_{n=1}^N$

Cost-Sensitive Label Embedding with Multidimensional Scaling

Training stage of CLEMS

- ▶ given training instances $\mathcal{D} = \{(\mathbf{x}^{(n)}, \mathbf{y}^{(n)})\}_{n=1}^N$ and cost function c
- ▶ determine the candidate set \mathcal{S}
- ▶ find $\mathbf{z}_i^{(t)}$ and $\mathbf{z}_i^{(p)}$ for all $\mathbf{y}_i \in \mathcal{S}$ by **multidimensional scaling**
- ▶ $\Phi: \mathbf{y}_i \rightarrow \mathbf{z}_i^{(p)}$
- ▶ train a multi-target regressor g from $\{(\mathbf{x}^{(n)}, \Phi(\mathbf{y}^{(n)}))\}_{n=1}^N$

Predicting stage of CLEMS

- ▶ given the testing instance (\mathbf{x}, \mathbf{y})
- ▶ $\Psi(\cdot) = \Phi^{-1}(\text{nearest neighbor}) = \Phi^{-1}(\text{argmin } d(\mathbf{z}_i^{(t)}, \cdot))$
- ▶ obtain the predicted embedded vector by $\tilde{\mathbf{z}} = g(\mathbf{x})$
- ▶ prediction $\tilde{\mathbf{y}} = \Psi(\tilde{\mathbf{z}})$

Experiments

Lists of experiments

- ▶ **comparison with label embedding algorithms**
 - ▶ LSDR algorithms ($\dim \mathcal{Z} < \dim \mathcal{Y}$)
 - ▶ LSDE algorithms ($\dim \mathcal{Z} \geq \dim \mathcal{Y}$)
- ▶ **comparison with cost-sensitive algorithms**
 - ▶ condensed filter tree (CFT) [Li et al., 2014]

Experiments

Lists of experiments


- ▶ **comparison with label embedding algorithms**
 - ▶ LSDR algorithms ($\dim \mathcal{Z} < \dim \mathcal{Y}$)
 - ▶ LSDE algorithms ($\dim \mathcal{Z} \geq \dim \mathcal{Y}$)
- ▶ **comparison with cost-sensitive algorithms**
 - ▶ condensed filter tree (CFT) [Li et al., 2014]

Settings


- ▶ 50% for training, 25% for validation, and 25% for testing
- ▶ base learner: random forest classifier or random forest regressor
- ▶ evaluation criteria
 - ▶ **F1 score** $\frac{2\|\mathbf{y} \cap \tilde{\mathbf{y}}\|_1}{\|\mathbf{y}\|_1 + \|\tilde{\mathbf{y}}\|_1}$ (\uparrow)
 - ▶ **Accuracy score** $\frac{\|\mathbf{y} \cap \tilde{\mathbf{y}}\|_1}{\|\mathbf{y} \cup \tilde{\mathbf{y}}\|_1}$ (\uparrow)
 - ▶ **Rank loss** $\sum_{\mathbf{y}[i] > \mathbf{y}[j]} (\mathbb{I}[\tilde{\mathbf{y}}[i] < \tilde{\mathbf{y}}[j]] + \frac{1}{2} \mathbb{I}[\tilde{\mathbf{y}}[i] = \tilde{\mathbf{y}}[j]])$ (\downarrow)
- ▶ average results of 20 experiments

Comparison with LSDR Algorithms

F1 score (\uparrow)




yeast




birds

Accuracy score (\uparrow)




yeast




birds

Rank loss (\downarrow)



yeast




birds


CLEMS outperforms LSDR algorithms on all the cost functions!

Comparison with LSDE Algorithms

F1 score (\uparrow)




yeast




birds

Accuracy score (\uparrow)




yeast




birds

Rank loss (\downarrow)



yeast



birds

CLEMS outperforms LSDE algorithms on all the cost functions!

Comparison with Cost-sensitive Algorithms

Table: Comparison with CFT

	F1 score (\uparrow)		Accuracy score (\uparrow)		Rank loss (\downarrow)	
	CFT	CLEMS	CFT	CLEMS	CFT	CLEMS
emo.	0.640	0.676	0.557	0.589	1.563	1.484
scene	0.703	0.770	0.656	0.760	0.723	0.672
yeast	0.649	0.671	0.543	0.568	8.566	8.302
birds	0.601	0.674	0.586	0.642	4.908	4.886
med.	0.635	0.814	0.613	0.786	5.811	5.170
enron	0.557	0.606	0.448	0.491	26.64	29.40
CAL.	0.371	0.419	0.237	0.273	1120.8	1247.9
EUR.	0.456	0.670	0.450	0.650	129.53	89.52

CLEMS outperforms CFT in most of the cases!

Conclusion

- ▶ **algorithm design:** cost-sensitive label embedding algorithm (CLEMS)
 - ▶ **mapping and nearest neighbor view** for the efficient decoding function
 - ▶ embed the cost information in **distance** by **multidimensional scaling**
 - ▶ **mirroring trick** for the asymmetric cost function
 - ▶ **candidate set** to reduce the computational burden
- ▶ **theoretical guarantee:**
 - ▶ prove the upper bound of the predicted cost for CLEMS
- ▶ **empirical performance:**
 - ▶ CLEMS outperforms the existing LSDR and LSDE algorithms
 - ▶ CLEMS is better than the state-of-the-art cost-sensitive algorithms

Conclusion

- ▶ **algorithm design:** cost-sensitive label embedding algorithm (CLEMS)
 - ▶ mapping and nearest neighbor view for the efficient decoding function
 - ▶ embed the cost information in distance by multidimensional scaling
 - ▶ mirroring trick for the asymmetric cost function
 - ▶ candidate set to reduce the computational burden
- ▶ **theoretical guarantee:**
 - ▶ prove the upper bound of the predicted cost for CLEMS
- ▶ **empirical performance:**
 - ▶ CLEMS outperforms the existing LSDR and LSDE algorithms
 - ▶ CLEMS is better than the state-of-the-art cost-sensitive algorithms

Thank you! Any question?